

Homework Assignment 0: Images as Functions

Due Monday, August 31st, 2020 at 11:59pm

Description

This problem set is really just to make sure you can load an image, manipulate the values, produce some output, and submit the code along with the report.

What to submit

Create a folder named `ps0_xxxx_LastName_FirstName`

- 1) `ps0_matlab_LastName_FirstName` if you are programming in MATLAB OR
- 2) `ps0_python_LastName_FirstName` if you are programming in Python

with the following structure and contents:

`ps0_xxxx_LastName_FirstName /`

- `input/` - directory containing input images, videos or other data supplied with the problem set
- `output/` - directory containing output images and other generated files

Note: Output images must be stored with following mandatory naming convention:

`ps<problem set #>-<question #>-<part>-<counter>.png`

Example: `ps0-1-a-1.png` (see question 1-a)

PNG format is easier to manipulate as it is lossless, hence we recommend using it over JPEG.

- `ps0.m` or `ps0.py` - your Matlab or Python code for this problem set
- `ps0_report.pdf` - A PDF file that shows all your output for the problem set, including images labeled appropriately (by filename, e.g. `ps0-1-a-1.png`) so it is clear which section they are for and the small number of written responses necessary to answer some of the questions (as indicated). Also, for each main section, if it is not obvious how to run your code please provide brief but clear instructions (no need to include your entire code in the report).
- `*.m` or `*.py` - Any other supporting files, including Matlab function files, Python modules, etc.

Zip it as `ps0_xxxx_LastName_FirstName.zip`, and submit on courseweb.

Questions

1. Input images

- a. Find two interesting images to use. They should be color, rectangular in shape (NOT square). Pick one that is wide and one tall.

You might find some classic vision examples [here](#). Or take your own. Make sure the image width or height do not exceed 512 pixels.

Output: Store the two images as ps0-1-a-1.png and ps0-1-a-2.png inside the output folder

2. Color planes

- a. Swap the red and blue pixels of image 1

Output: Store as ps0-2-a-1.png in the output folder

- b. Create a monochrome image (img1_green) by selecting the green channel of image 1

Output: ps0-2-b-1.png

- c. Create a monochrome image (img1_red) by selecting the red channel of image 1

Output: ps0-2-c-1.png

- d. Which looks more like what you'd expect a monochrome image to look like? Would you expect a computer vision algorithm to work on one better than the other?

Output: Text response in report ps0_report.pdf

3. Replacement of pixels (Note: For this, use the *better* channel from 2-b/2-c as monochrome versions.)

- a. Take the center square region of 100x100 pixels of monochrome version of image 1 and insert them into the center of monochrome version of image 2

Output: Store the new image created as ps0-3-a-1.png

4. Arithmetic and Geometric operations

- a. What is the min and max of the pixel values of img1_green? What is the mean? What is the standard deviation? And how did you compute these?

Output: Text response, with code snippets

- b. Compute and plot the histogram of img1_green.

Output: ps0-4-b-1.png, text response: What is image histogram?

- c. Subtract the mean from all pixels, then divide by standard deviation, then multiply by 10 (if your image is 0 to 255) or by 0.05 (if your image ranges from 0.0 to 1.0). Now add the mean back in.

Output: ps0-4-c-1.png

- d. Now, compute and plot the histogram of the resultant image in part c.

Output: ps0-4-d-1.png, text response: compare between the histograms you obtained in b and d.

- e. Shift img1_green to the left by 2 pixels.

Output: ps0-4-e-1.png

- f. Subtract the shifted version of img1_green from the original, and save the difference image.

Output: ps0-4-f-1.png (make sure that the values are legal when you write the image so that you can see all relative differences), text response: What do negative pixel values mean anyways?

5. Noise

- a. Take the original colored image (image 1) and start adding Gaussian noise to the pixels in the green channel. Increase sigma until the noise is somewhat visible.
Output: ps0-5-a-1.png, text response: What is the value of sigma you had to use?
- b. Compute and plot the histogram of the noisy image.
Output: ps0-5-b-1.png, text response: how does it compare to the histogram from part 4.b?
- c. Now, instead add that amount of noise to the blue channel.
Output: ps0-5-c-1.png
- d. Which looks better? Why?
Output: Text response
- e. Filter the noisy image in part 5.a. using a median filter (e.g., `medfilt2` in MATLAB), show the output image and its histogram.
Output: ps0-5-e-1.png (filtered output image), and ps0-5-e-2.png (histogram of the output image)
Text response: compare the output image and its histogram to those of the noisy image.
- f. Repeat e, but use a Gaussian filter (e.g., `imgaussfilt` in MATLAB) instead
Output: ps0-5-f-1.png (filtered output image), and ps0-5-f-2.png (histogram of the output image)
Text response: compare the output image and its histogram to those of the noisy image. Also compare between the outputs from parts c and d. Which filter is better?